

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

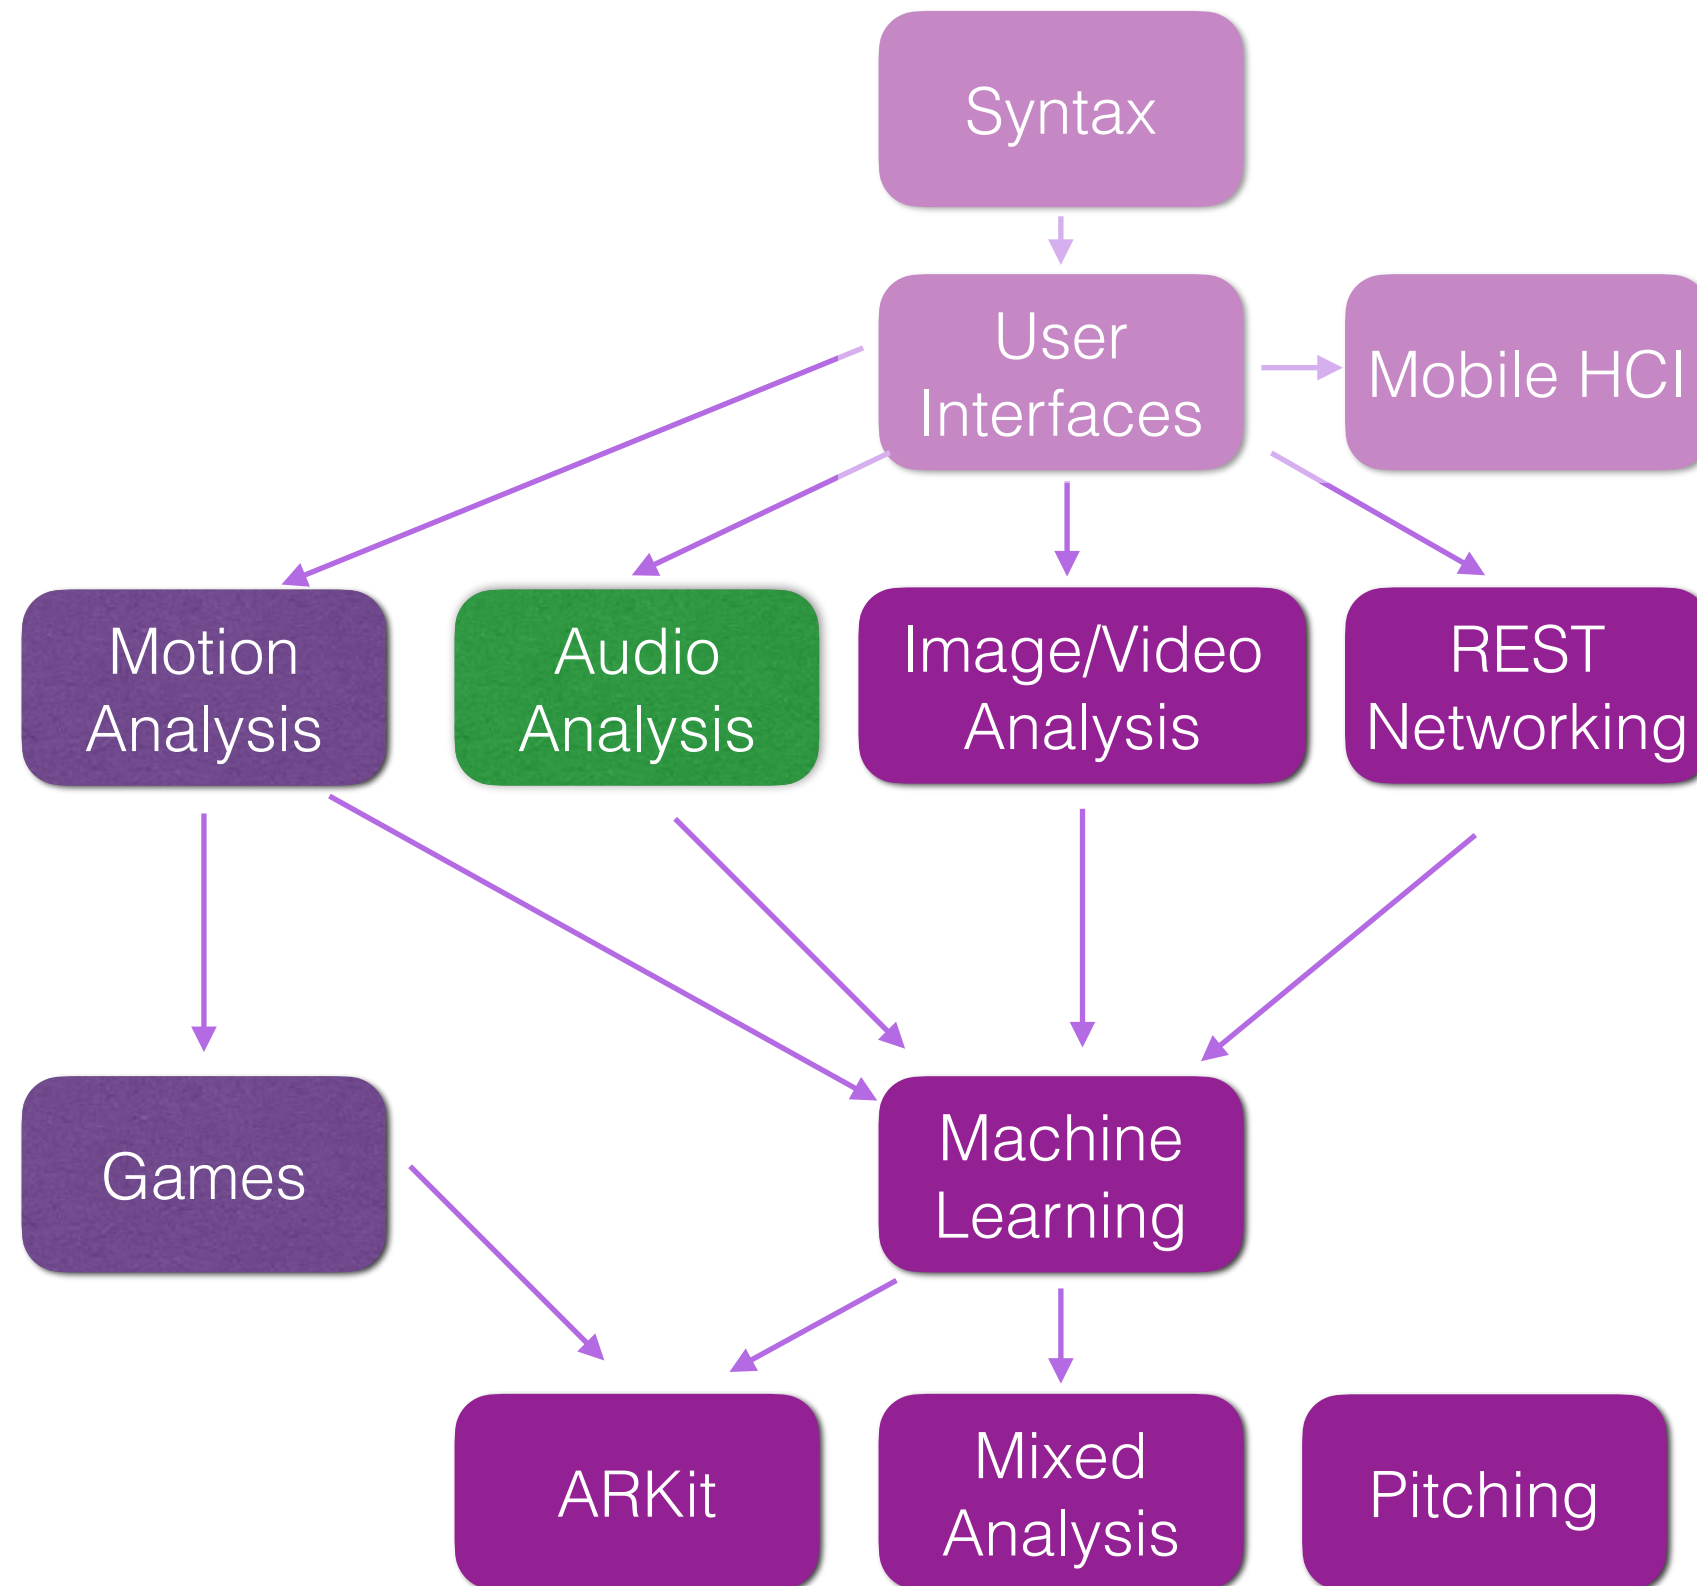
audio graphing, sampled data, & accelerate

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

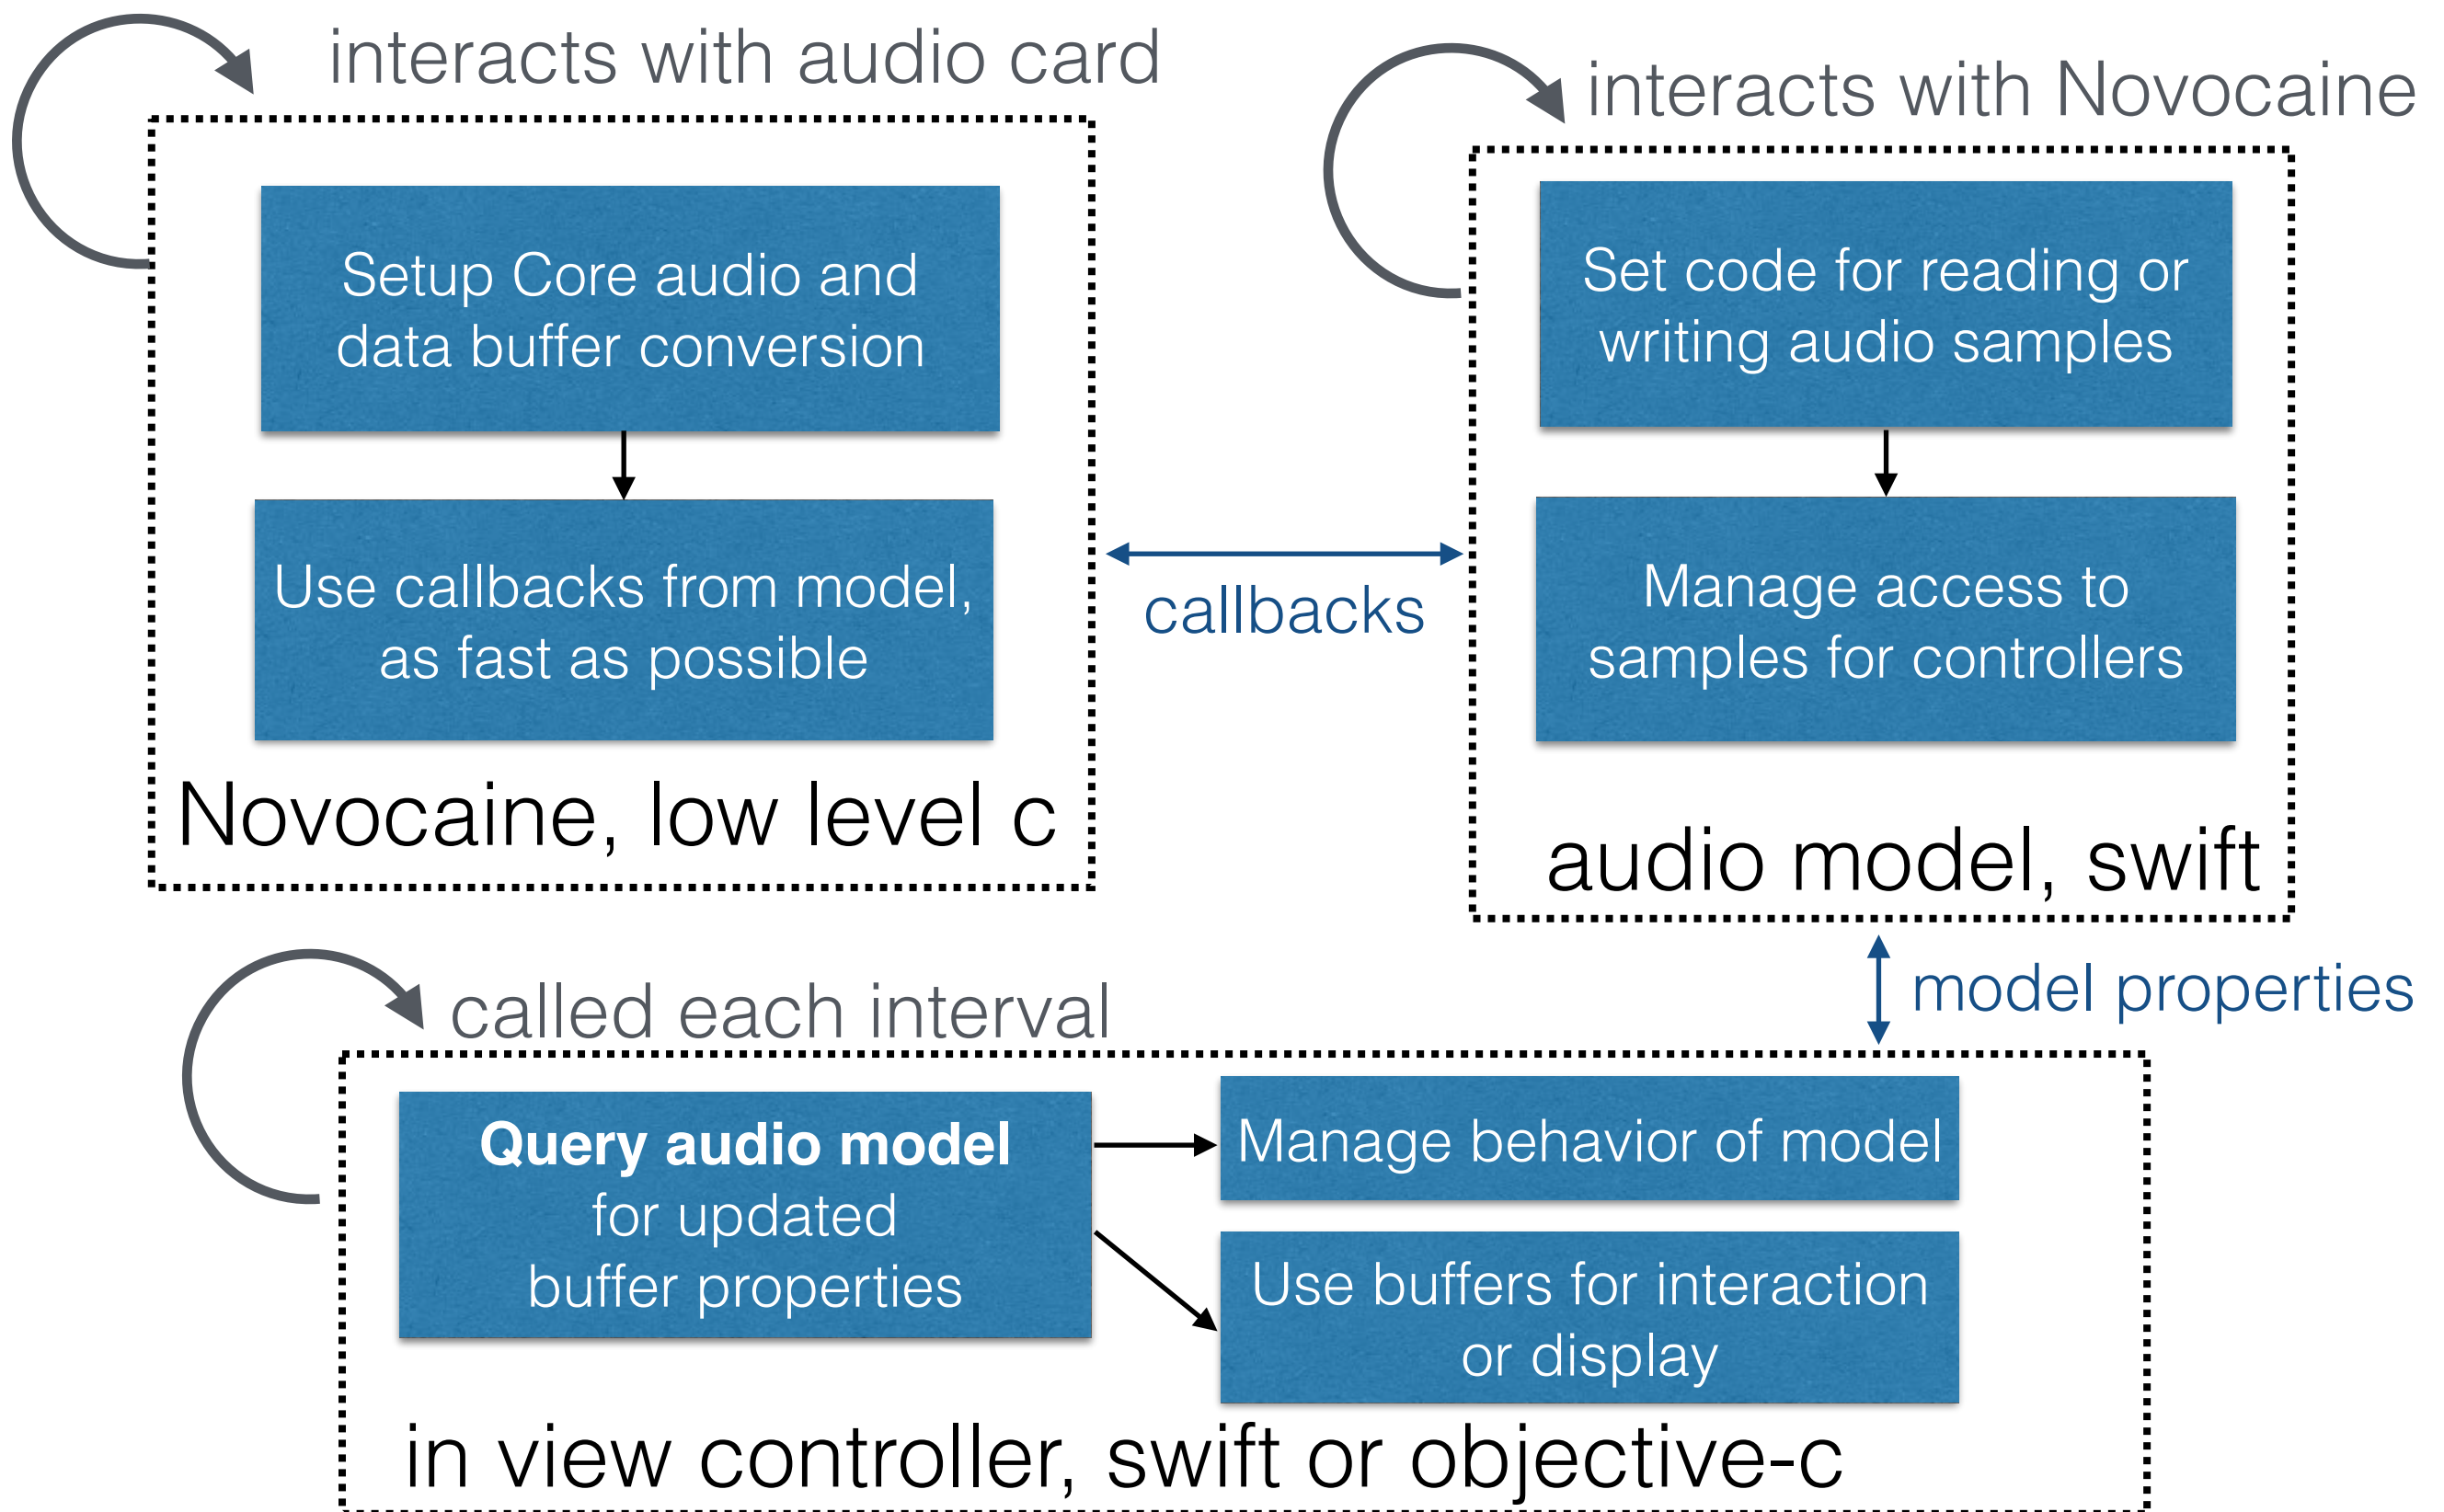
agenda and logistics

- logistics
 - **flipped module on audio next time!**
 - TA Office Hours: See canvas
- agenda
 - dealing with sampled data
 - the accelerate framework
 - massive digital signal processing library
 - graphing audio fast (well, graphing anything)
 - must use lowest level graphing, Metal

class overview



review: MVC with audio



sample from the mic

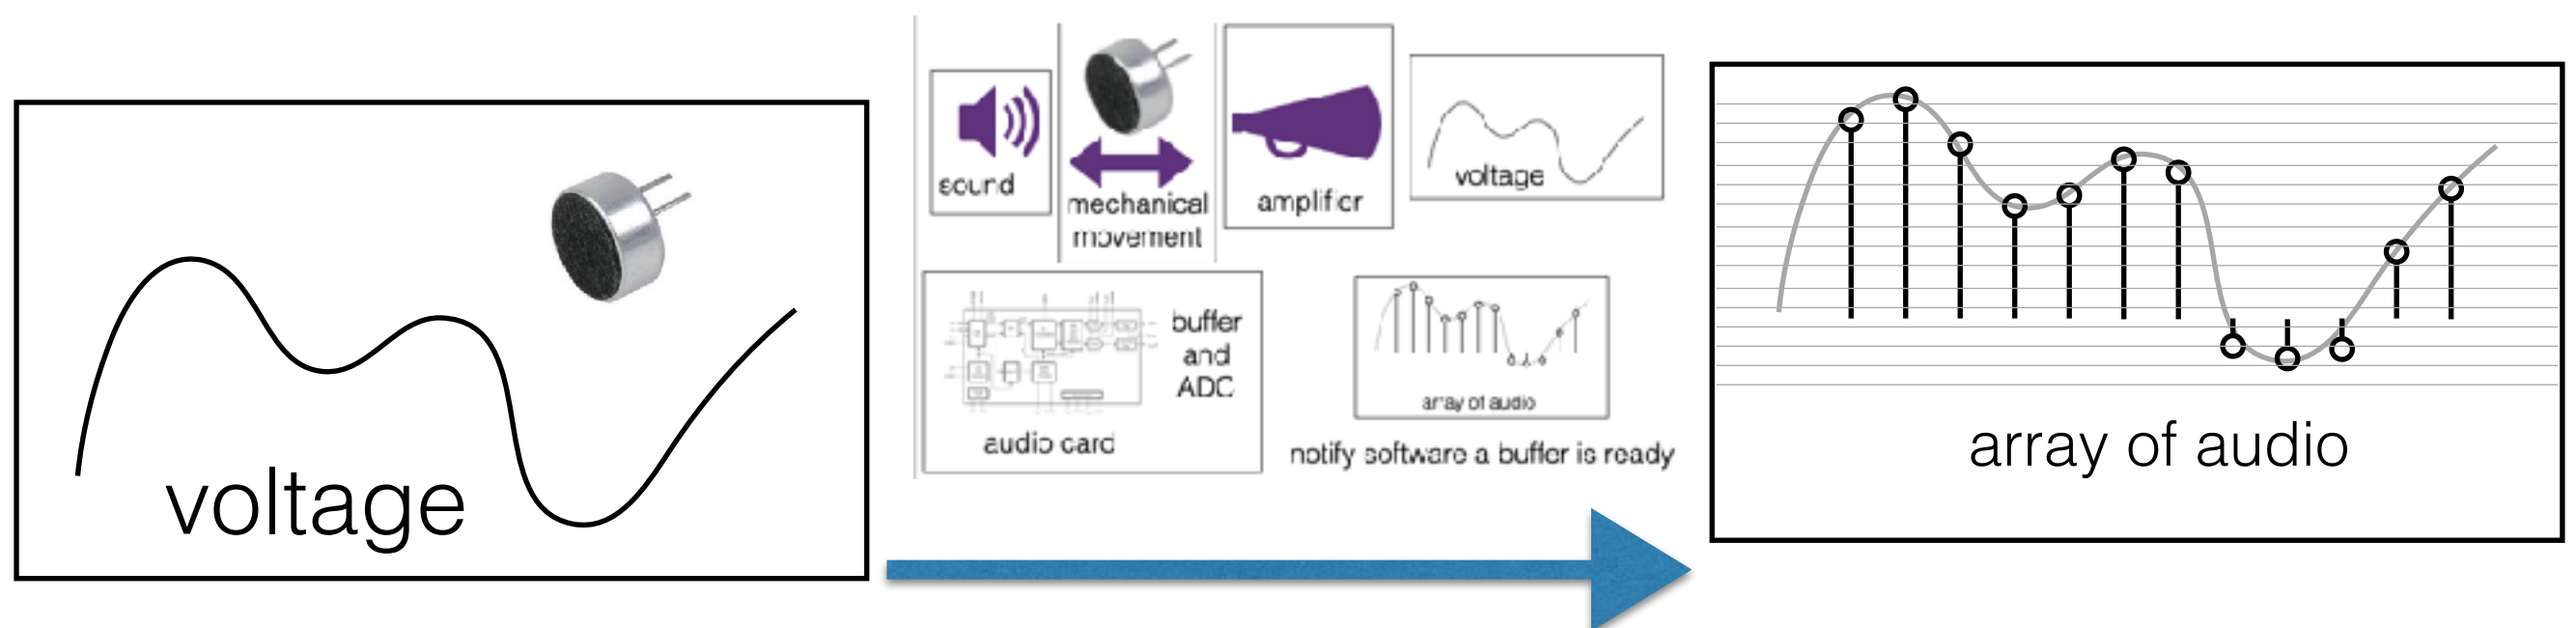
- recall: data from the microphone on novocaine



```
private func handleMicrophone (data:Optional<UnsafeMutablePointer<Float>>,
                                numFrames:UInt32,
                                numChannels: UInt32) {
    if let arrayData = data{
        // 🎤 -> 📈 get max element in the buffer
        // bonus: vDSP example (will cover in next lecture)
        // example with iOS accelerate to quickly handle the array
        var max:Float = 0
        vDSP_maxv(arrayData, 1, &max, vDSP_Length(numFrames))
        print(max)
    }
}
```


intro to sampled data

- physical processes are continuous
 - digitization **may change** how we **understand** the signal
- digitization occurs in time and amplitude
 - time: sampling
 - amplitude: quantization



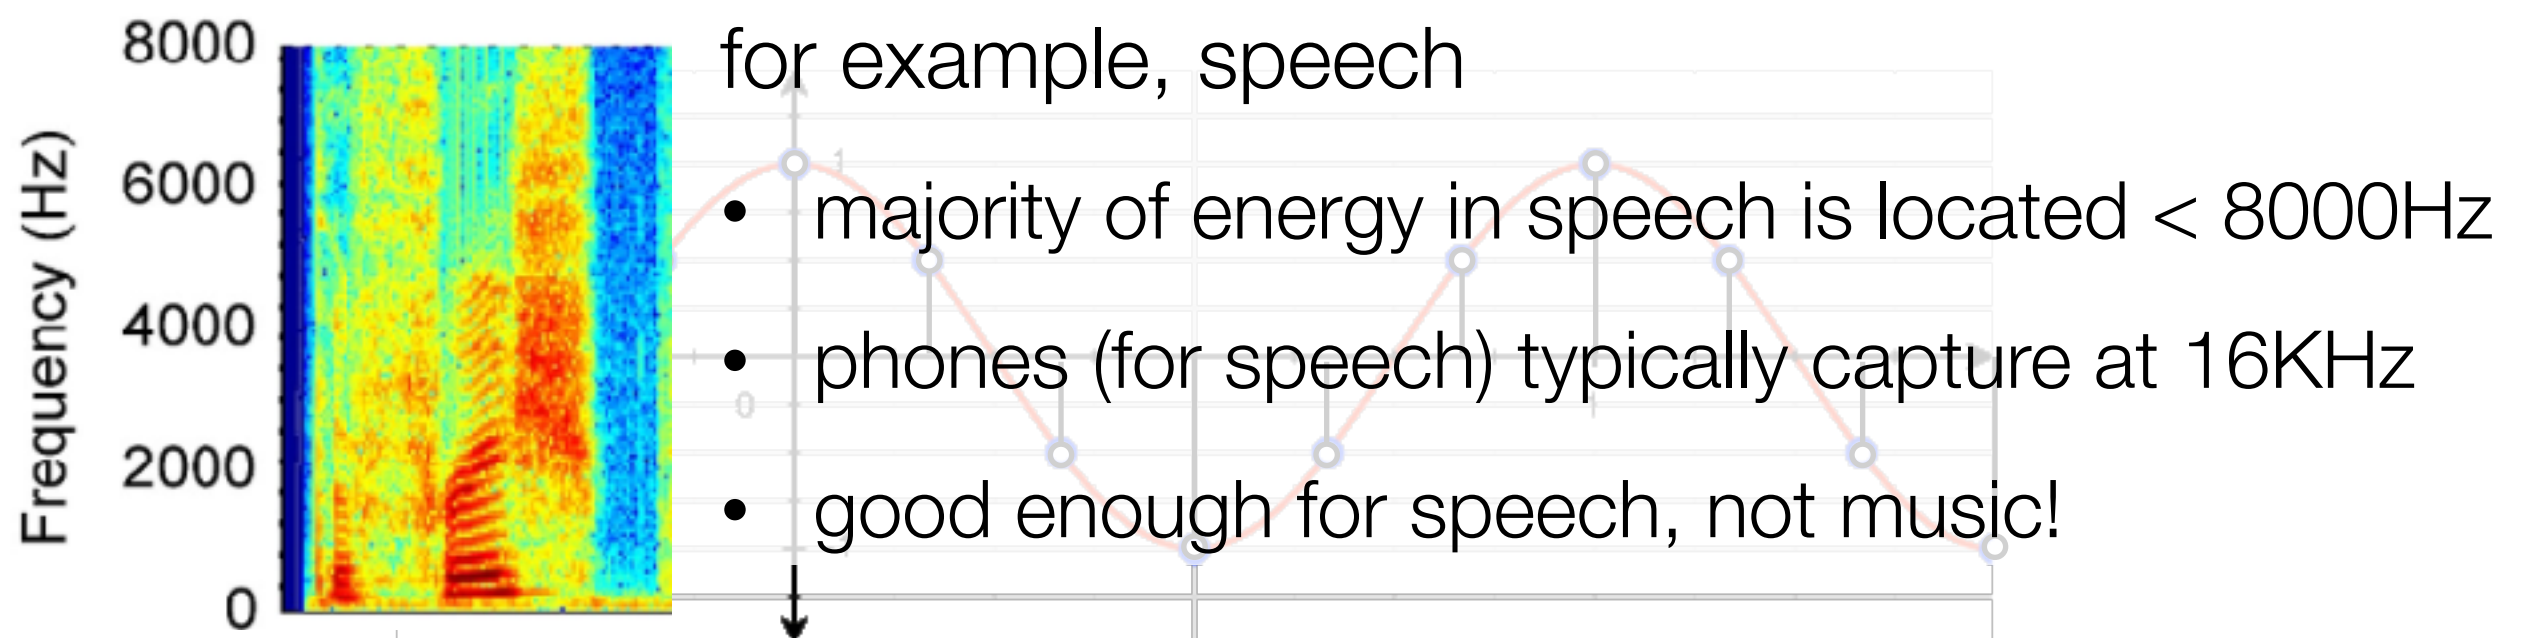
sampled data

- quantization (amplitude)
 - introduces error in estimating amplitude of a signal
 - error can be reduced by adding more “bits per sample”
- most ADCs are 16 bits, considered “good enough”
- sufficient for most uses
 - not for others!



sampling errors

- sampling in time
 - introduces errors through ‘aliasing’, limits the range of frequencies able to be accurately captured
- heuristics
 - **Nyquist:** if capturing an “ F ”Hz signal, need to sample at least $2 \times “F”$ Hz
 - changing sample rates is complicated



sanity check

- I need to detect an 80Hz signal
 - what sampling rate should we use?
- I want to detect a feather dropping next to the microphone
 - can the sound be detected?

making a sine wave

- we want to create a sine wave and play it to the speakers


$$g(t) = \sin(2\pi f t) \quad \text{equation for sine wave}$$

f , frequency in Hz

t , time in “seconds”

but we are working digitally, so we have an “index” in an array,
not time!

Output Array in Novocaine `g:Optional<UnsafeMutablePointer<Float>>`



making a sine wave

$$g[n] = \sin \left(2\pi f \left(\frac{n}{F_s} \right) \right) \quad \text{how to program this?}$$

```
for (int n=0; n < numFrames; ++n)
{
    data[n] = sin(2*M_PI*frequency*n/samplingRate);
}
```

is this efficient?

should this be initialized inside the audio callback?


```
float phase = 0.0;
double phaseIncrement = 2*M_PI*frequency/samplingRate;
for (int n=0; n < numFrames; ++n)
{
    data[n] = sin(phase);
    phase += phaseIncrement;
}
```

$$g[n] = \sin(\theta)$$

$$\theta \leftarrow \theta + \frac{2\pi f}{F_s}$$


sine wave discontinuity

$$g[n] = \sin\left(2\pi f \left(\frac{n}{F_s}\right)\right) \quad \theta \leftarrow \theta + \frac{2\pi f}{F_s} \quad g[n] = \sin(\theta)$$



phase = 0 in each output block

```
float phase = 0.0;
double phaseIncrement = 2*M_PI*freq/samplingRate;
for (int n=0; n < numFrames; ++n)
{
    data[n] = sin(phase);
    phase += phaseIncrement;
}
```



phase variable
overflows,
discontinuity!

making a sine wave

- bringing it all together




$$\theta \leftarrow \theta + \frac{2\pi f}{F_s} \quad g[n] = \sin(\theta)$$

```
var frequency = 18000.0 //starting frequency
var phase = 0.0
var samplingRate = audioManager.samplingRate
let sineWaveRepeatMax = 2*Double.pi
```

```
outputBlockFunction(data:(...), numFrames:(UInt32), numChannels:(UInt32))
{
    var phaseIncrement = 2*Double.pi*frequency/samplingRate
    var i=0
    while (i < numFrames)
    {
        data[i] = sin(phase)
        i += 1
        phase += phaseIncrement
        if (phase >= sineWaveRepeatMax) { phase -= sineWaveRepeatMax }
    }
}
```

what if more than one channel?


what if the frequency changes?



phase
overflow

play samples to speakers

- demo, play sine wave



the accelerate framework

- very powerful digital signal processing (DSP) library
 - look at vDSP Programming Guide on developer.apple.com for the complete API
- provides mathematics for performing fast DSP
 - fast is the name of the game — it uses “c”
- SIMD Single Instruction, Multiple Data

input data

stride

scalar

output

array length

```
vDSP_vsmul(data, 1, &mult, data, 1, numFrames*numChannels);
```

```
void vDSP_vsmul (  
    const float __vDSP_input1[],  
    vDSP_Stride __vDSP_stride1,  
    const float *__vDSP_input2,  
    float __vDSP_result[],  
    vDSP_Stride __vDSP_strideResult,  
    vDSP_Length __vDSP_size  
);
```

https://developer.apple.com/documentation/accelerate/1450020-vdsp_vsmul

examples

what do each of these implement?


```
outputBlockFunction(data:(...), numFrames:(UInt32), numChannels:(UInt32)) {  
    ringBuffer.fetchFreshData(data, withNumFrames:numFrames)  
    var volume = userSetMultiplyFromSlider;  
    vDSP_vsmul(data, 1, &volume, data, 1, numFrames*numChannels)  
}
```

```
inputBlockFunction(data:(...), numFrames:(UInt32), numChannels:(UInt32)) {  
    // get the max  
    var maxVal = 0.0;  
    vDSP_maxv(data, 1, &maxVal, numFrames*numChannels);  
  
    print("Max Audio Value: %f\n", maxVal);  
}
```

```
inputBlockFunction(data:(...), numFrames:(UInt32), numChannels:(UInt32)) {  
    vDSP_vsqr(data, 1, data, 1, numFrames*numChannels);  
    var meanVal = 0.0;  
    vDSP_meanv(data, 1, &meanVal, numFrames*numChannels);  
}
```


audio graphing

- we want to see the incoming samples
- good for debugging
- equalizers, oscilloscope type applications, etc.



how much data to show?

- sampling at 48kHz == 48000 samples per second



graph 0.5 second
window is:
24000 samples

display is >640
pixels wide

what if we want
lots of graphs?

solution

- use the GPU
- set vectors of data on a 2D plane
- let the renderer perform scaling, anti-aliasing, and bit blitting to screen
- ...this is not a graphics course
- ...but we need to use the Metal API


Metal



Apple used the mobile [multiplayer online battle arena](#) game [Vainglory](#) to demonstrate Metal's graphics capabilities at the [iPhone 6's](#) September 2014 announcement event^[1]

Developer(s)	Apple Inc.
Initial release	June 2014; 6 years ago
Stable release	3 / June 2019; 1 year ago
Written in	Shading Language: C++14 , Runtime/API: Objective-C
Operating system	iOS, iPadOS, macOS, tvOS
Type	3D graphics and compute API
License	proprietary
Website	developer.apple.com/metal/

the MetalGraph class



drag class/shaders
into project, if needed

```
lazy var graph: MetalGraph? = {  
    return MetalGraph(mainView: self.view)  
}()
```

declare and init property

```
// add in a graph for displaying the audio  
graph.addGraph(withName: "time",  
               shouldNormalize: false,  
               numPointsInGraph: AUDIO_BUFFER_SIZE)
```

add graph names to controller
and how many expected points in array

```
// periodically, display the audio data  
graph.updateGraph(  
    data: timeData,  
    forKey: "time"  
)
```

refresh data for each
named graph key

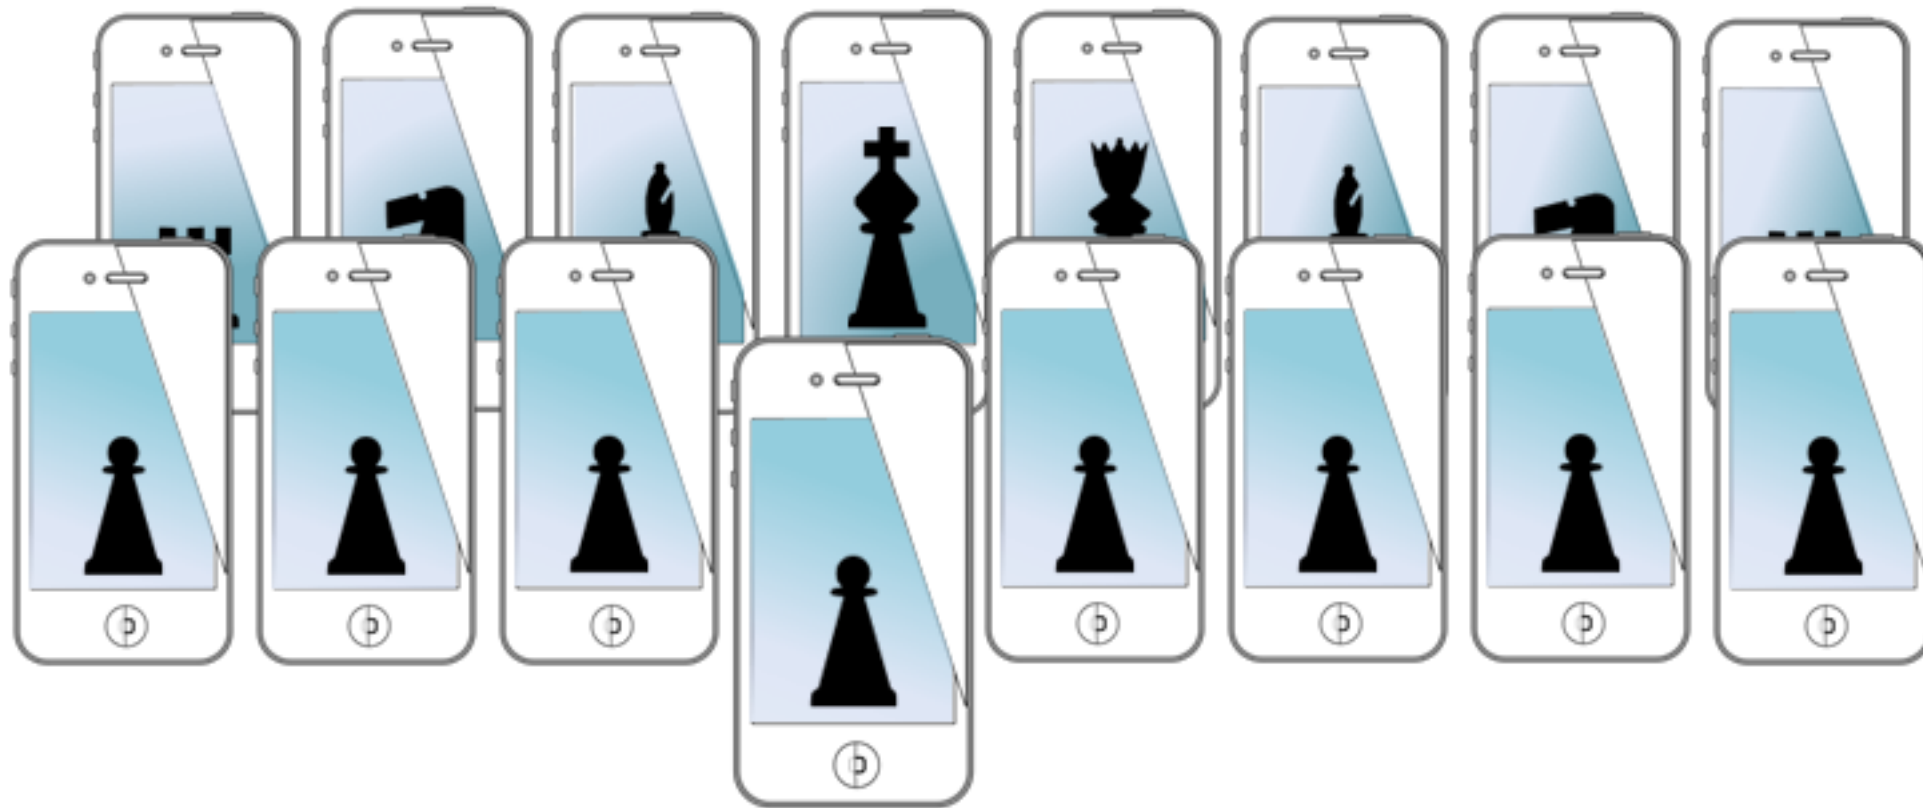
Properties: automatic screensize (pixel) downsampling
automatic coloring based in iOS scheme,
efficient memory management through vertex buffers
normalize: (default) assume **data is between -1.0 to 1.0**

audio graphing demo!



if time: (1) add another graph
(2) preview of FFT graphing

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

audio graphing, sampled data, & accelerate

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

Supplemental Slides: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

Supplemental Slides

- these slides were removed from the course because of their complexity
- you need a good background in signal representation and time series analysis to really understand these slides

Optional Concepts not Covered in Course Anymore

filters!

- we will cover what we can...

signals and systems

- signals are collections of sampled data (arrays)
 - such as audio, accelerometer, etc.
 - can also be 2D, like images
- systems are objects which manipulate signals
 - characterized by their “input/output” relationships
 - we say “ $x[n]$ is passed through H , resulting in $y[n]$ ”

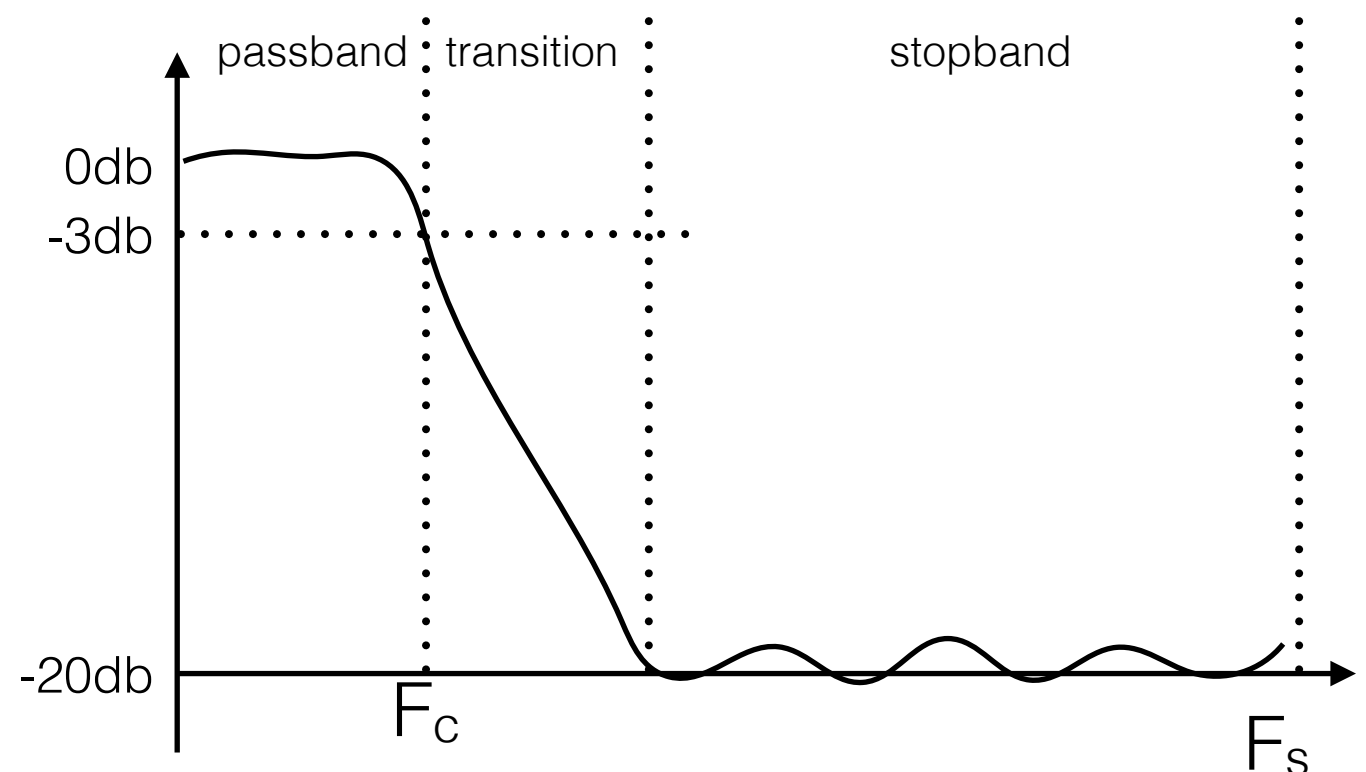


filters

- filters are systems which manipulate frequencies
 - certain frequencies to pass through, but not others
 - “lowpass” filter allows low frequencies to pass through
 - “highpass” filter likewise allows high frequencies through
- keep in mind: no filter is perfect!!
 - no filter will pass everything you want while stopping everything you don't
 - everything is a balance between different parameters you can control
- we won't study how to design filters
 - we will study properties of filters and how to use them
 - so we need to know what filters can and cannot do

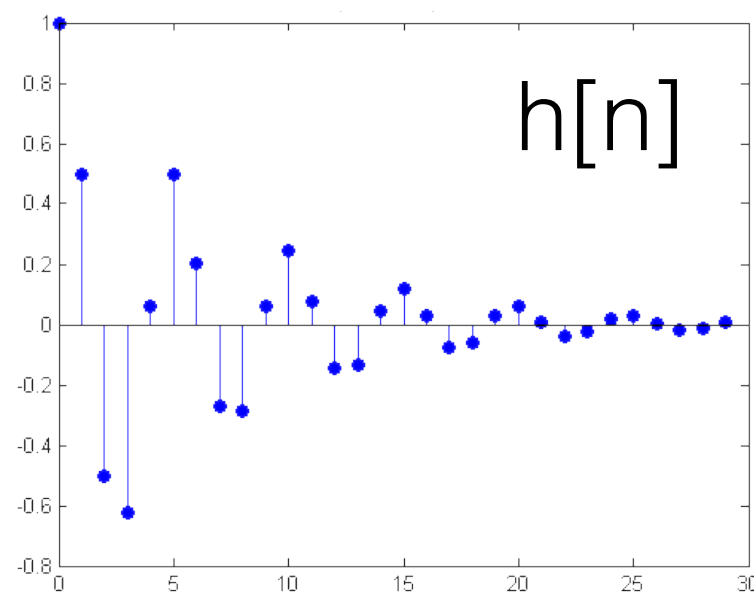
filters in frequency

- filters can be characterized a few different ways
 - let's start by looking at their properties in the frequency domain
- filters have the following frequency-domain attributes:
 - passband gain
 - passband bandwidth
 - stopband attenuation
 - transition bandwidth



filters in time

- filter are also signals (time series)
 - the series is called the “impulse response” of the filter
 - the frequency-domain plots are just Fourier transforms of the impulse response (magnitude)
- the time-domain property we care about is length
 - everything else is best left to a filter design course



so how to design a filter?

- scipy.signal in python (try to use remez)
- decent tutorial:
 - <http://mpastell.com/2010/01/18/fir-with-scipy/>
- matlab
 - fdatool
- lots of other places!


```
from pylab import *  
import scipy.signal as signal  
n = 61  
a = signal.firwin(n, cutoff = 0.3,  
                  window = "hamming")
```

filtering by convolution



- we apply a filter using **convolution**
 - convolution allows us to combine frequency properties of two signals without taking an FFT
- basic principle:
 - convolution in time is multiplication in frequency
 - so the filter's frequency response will be multiplied by the frequency response of the signal

$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$

convolution



$$y[n] = \sum_{i=0}^{N-1} h[n-i]x[i]$$



length


N

M

$N+M-1$

convolution efficiency


- algorithmic complexity
 - convolution is not particularly efficient, $O(N \times M)$
- to convolve faster, use that Fourier property:
 - “convolution in time is multiplication in frequency”
- why not just multiply to begin with!



its circular



- just using N point FFT performs “Circular Convolution”
 - which is not linear convolution
 - causes the tail end of the convolution to “wrap around” to the beginning
 - FFT assumes the function is periodic (we did not talk about this)
- be aware of circularity when filtering your signal with the FFT
 - zero-padding can solve this for you!
 - zero-pad both signals to a length that will contain the entire convolution, $N+M-1$
 - for streaming, you must use overlap-and-add!
 - http://en.wikipedia.org/wiki/Overlap-add_method

overlap and add







windowing: spectral BW widening





- multiplication in time is convolution in frequency
- a window is something we multiply in time with our signal
- windowing is unavoidable
 - why? we cannot take an infinite FFT...



windowing: spectral leakage



windowing: spectral leakage



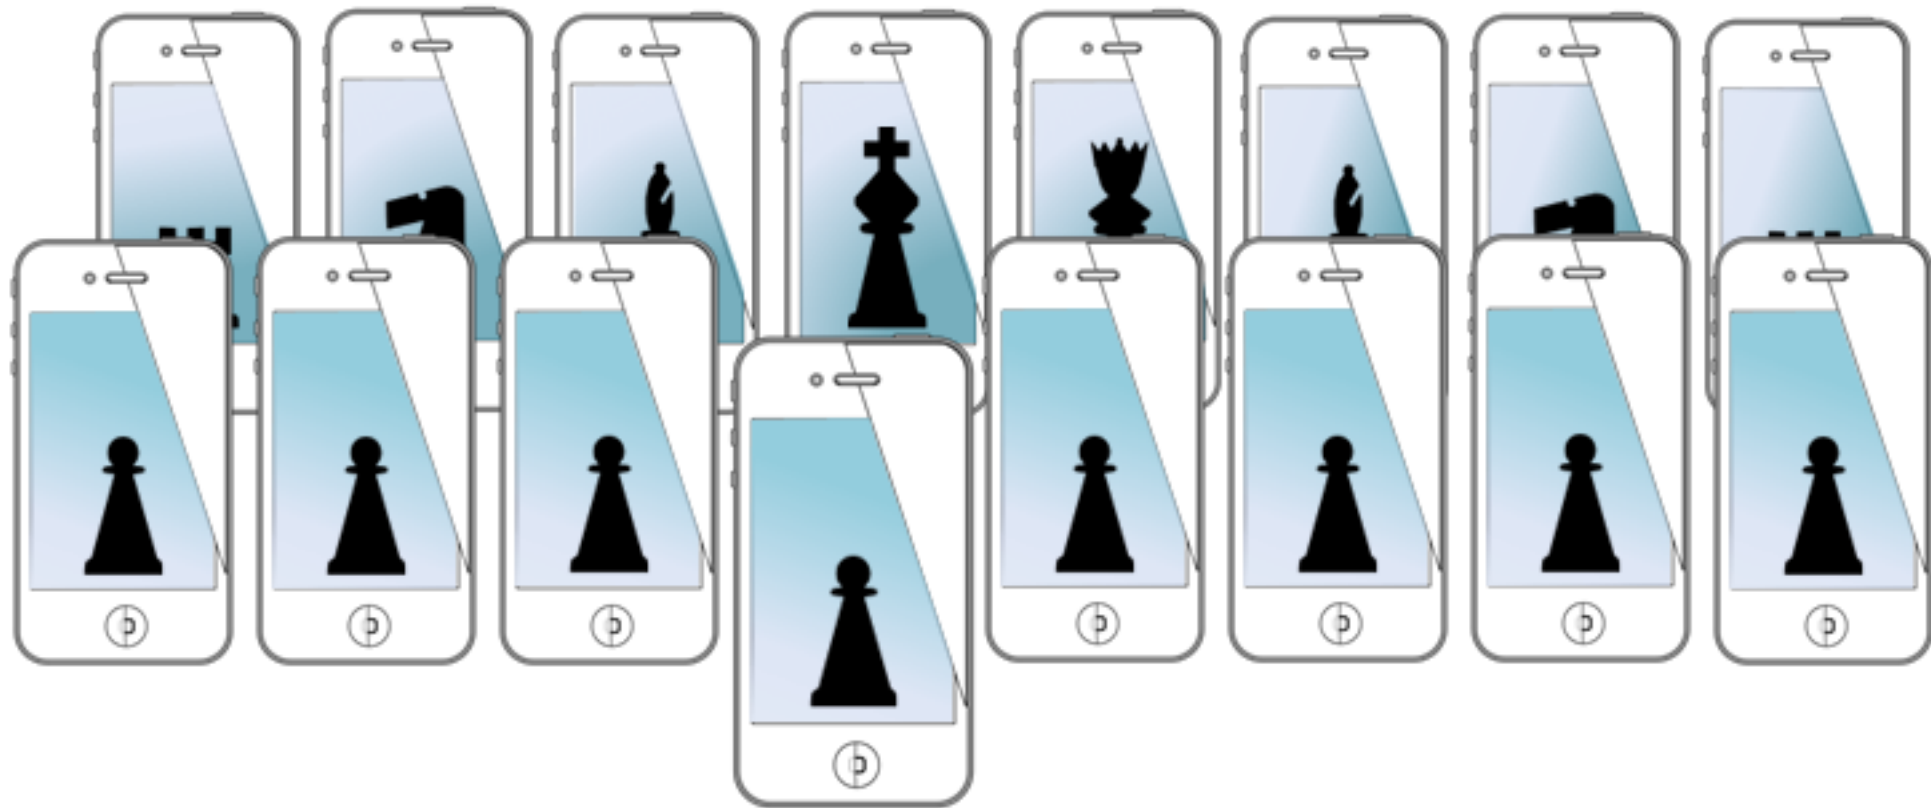
which window to use?

- depends
 - narrowest main lobe: rect
 - good tradeoff: hamming (or Von Hann)
 - optimal tradeoff for a given bandwidth:
 - discrete prolate spheroidal sequence (dpss, Slepian taper)

FFT review

- sampling rate
 - dictates the time between each sample, $(1 / \text{sampling rate})$
 - max frequency we can measure is half of sampling rate
- resolution in frequency
 - tradeoff between length of FFT and sampling rate
 - each frequency “bin” is an index in the FFT array
 - each bin represents (F_s / N) Hz
 - what does that mean for 12 Hz accuracy?
- windowing is a result of “convolution” in frequency
 - some windows prevent “leakage” at the cost of frequency resolution

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

Supplemental Slides: filtering and windowing

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University