

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

Speech Recognition and Dictation

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University

course logistics and agenda

- logistics:
 - final projects coming soon!!
- agenda:
 - speech recognition (from Speech Framework, SF)
 - SFSpeechRecognizer
 - **I helped Apple wreck a nice beach!**

overview

- introduced in 2016, same technology underpinning Siri
 - user must provide explicit authorization
 - free, but limited to certain number of recognitions per day
 - only allowed to dictate about 1 minute of audio
 - supports streaming audio and file I/O



general usage

- uses API similar to REST (like Core Vision)
 - create a task
 - configure it (options)
 - start task
 - use completion handler (for updates and final text)
- best practices:
 - signify to user that the app is recording
 - show dictation as it happens

tradeoff server/on-device

- 2016: speech is translated via cloud services
- 2019: also available on device

	Server	On-device
Accuracy	Best	Good
Limits	1 minute max audio duration Limited requests per day	None
Languages	50+	10+

using SFSpeechRecognizer



- callback model
- needs AVFoundation for adding audio chunks

```
private let speechRecognizer = SFSpeechRecognizer(locale: Locale(identifier: "en-US"))!  
private var recognitionRequest: SFSpeechAudioBufferRecognitionRequest?  
private var recognitionTask: SFSpeechRecognitionTask?  
private let audioEngine = AVAudioEngine()  
  
let audioSession = AVAudioSession.sharedInstance()  
audioSession.setCategory(AVAudioSession.Category.record)  
audioSession.setMode(AVAudioSession.Mode.measurement)  
audioSession.setActive(true, options: .notifyOthersOnDeactivation)  
  
let inputNode = audioEngine.inputNode  
let recordingFormat = inputNode.outputFormat(forBus: 0)  
  
inputNode.installTap(onBus: 0, bufferSize: 1024, format: recordingFormat)  
{ (buffer: AVAudioPCMBuffer, when: AVAudioTime) in  
    self.recognitionRequest?.append(buffer)  
}  
  
audioEngine.prepare()  
audioEngine.start()
```

much of the code also has **guards** for **error checking**

using SFSpeechRecognizer



```
let inputNode = audioEngine.inputNode
let recordingFormat = inputNode.outputFormat(forBus: 0)

inputNode.installTap(onBus: 0, bufferSize: 1024, format: recordingFormat)
{ (buffer: AVAudioPCMBuffer, when: AVAudioTime) in
    self.recognitionRequest?.append(buffer)
}

// perform on device, if possible
if speechRecognizer.supportsOnDeviceRecognition {
    recognitionRequest?.requiresOnDeviceRecognition = true
}

recognitionTask = speechRecognizer.recognitionTask(with: recognitionRequest)
{ [unowned self] result, error in
    if let result = result {
        let transcribedText = result.bestTranscription.formattedString
        // do something with text
    }

    if result?.isFinal ?? (error != nil) {
        // this will remove the listening tap
        // so that the transcription stops
        inputNode.removeTap(onBus: 0)
    }
}
```

more advanced speech processing

- recognition result contains many aspects of the voice, including:
 - Transcribed text (as we have seen)
 - Alternate transcriptions
 - Confidence in result
 - Timing
 - Speaking rate
 - Pause duration
 - Voice analytics

more advanced speech processing

```
bestTranscription=<SFTtranscription>,
formattedString=I helped Apple recognize speech,

speakingRate=0.000000, averagePauseDuration=0.000000,

segments=(
  <SFTtranscriptionSegment>, substringRange={0, 1}, timestamp=0.54, duration=0.24,
  confidence=0.966,
  substring=I, alternativeSubstrings=(\n),
  phoneSequence=AY,
  ipaPhoneSequence=\U02c8a\U0361\U026a, voiceAnalytics=(null),

  <SFTtranscriptionSegment>, substringRange={2, 6}, timestamp=0.78,
  duration=0.36000000000000001, confidence=0.966,
  substring=helped, alternativeSubstrings=(\n),
  phoneSequence=h EH l p t,
  ipaPhoneSequence=h.\U02c8\U025b.l.p.t,
  voiceAnalytics=(null),

  ...

  <SFTtranscriptionSegment>, substringRange={25, 31}, timestamp=2.49,
  duration=0.56999999999999998, confidence=0.966,
  substring=speech, alternativeSubstrings=(\n),
  phoneSequence=s p EE ch,
  ipaPhoneSequence=s.p.\U02c8i.t\U0361\U0283, voiceAnalytics=(null)
),
```

running on device will limit the available features!

more advanced speech processing

Each segment now has incredible amount of information

```
<SFTranscriptionSegment>,  
substringRange={0, 10}, timestamp=0.27, duration=0.65, confidence=0.911,  
substring=Performing,  
alternativeSubstrings=(\n),  
phoneSequence=(null),  
ipaPhoneSequence=(null),
```

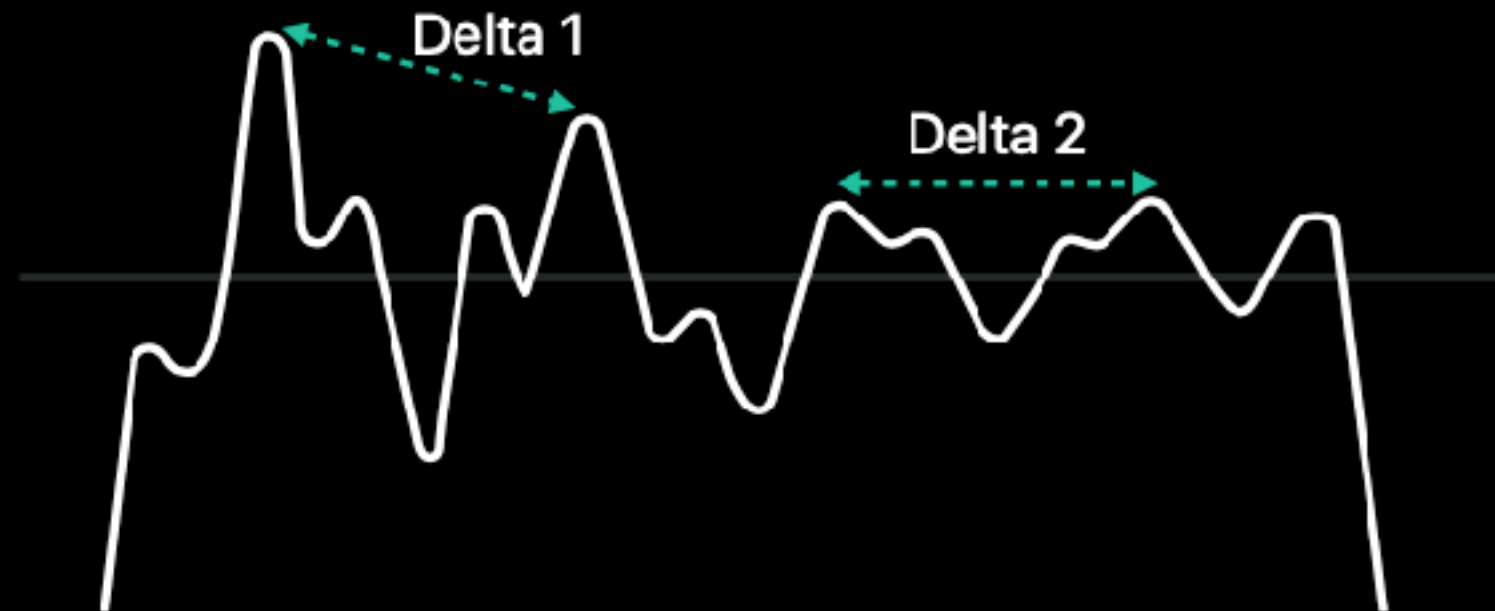
```
voiceAnalytics=<SFVoiceAnalytics>,  
jitter=<SFAcousticFeature>, featureValues=(12.53122 ... 0.6218916),  
frameDuration=0.010000,  
shimmer=<SFAcousticFeature>, featureValues=(0.7158176 ... 2.518468),  
frameDuration=0.010000,  
pitch=<SFAcousticFeature>, featureValues=(0.8526305, ..., 0.04258926),  
frameDuration=0.010000,  
voicing=<SFAcousticFeature>, featureValues=(0.07444749 ... 0.4056852),  
frameDuration=0.010000",
```

each featureValues array is length of
audio frames, could be **hundreds** of values

analytics

Jitter

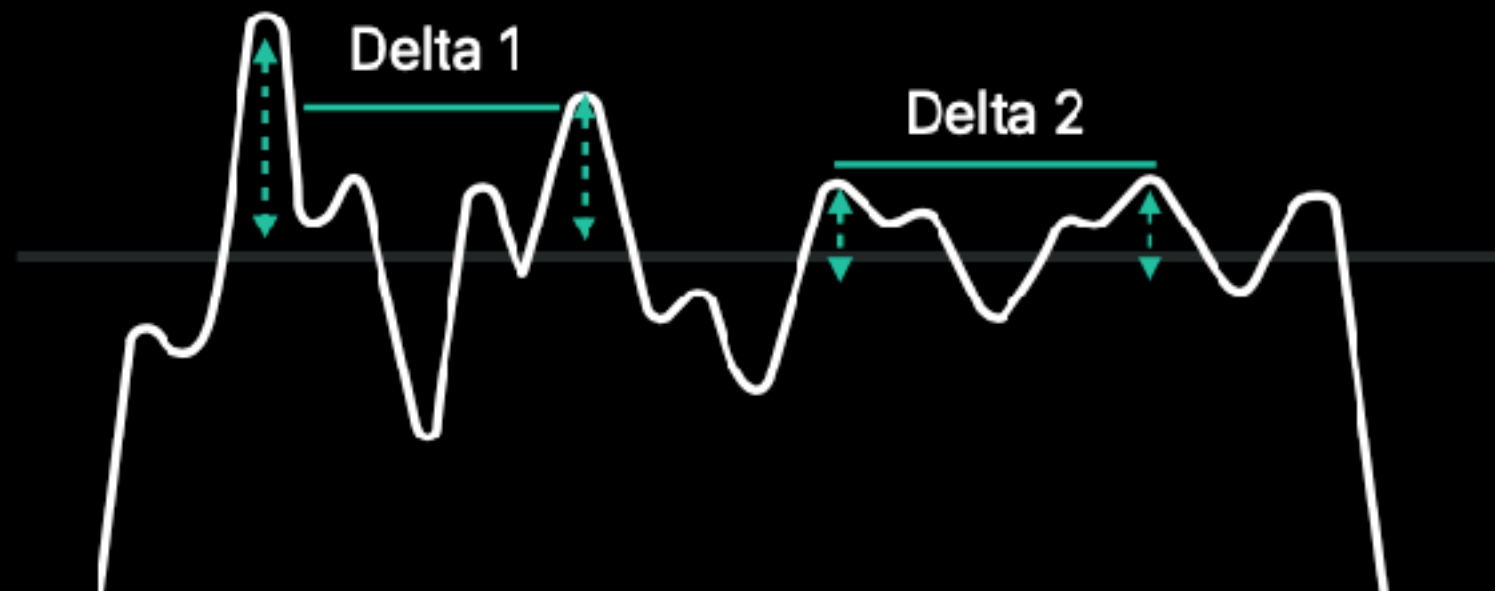
Measures variation in pitch



$$\text{Jitter} = \frac{\Delta 1 - \Delta 2}{\text{mean}}$$

Shimmer

Measures variation in amplitude



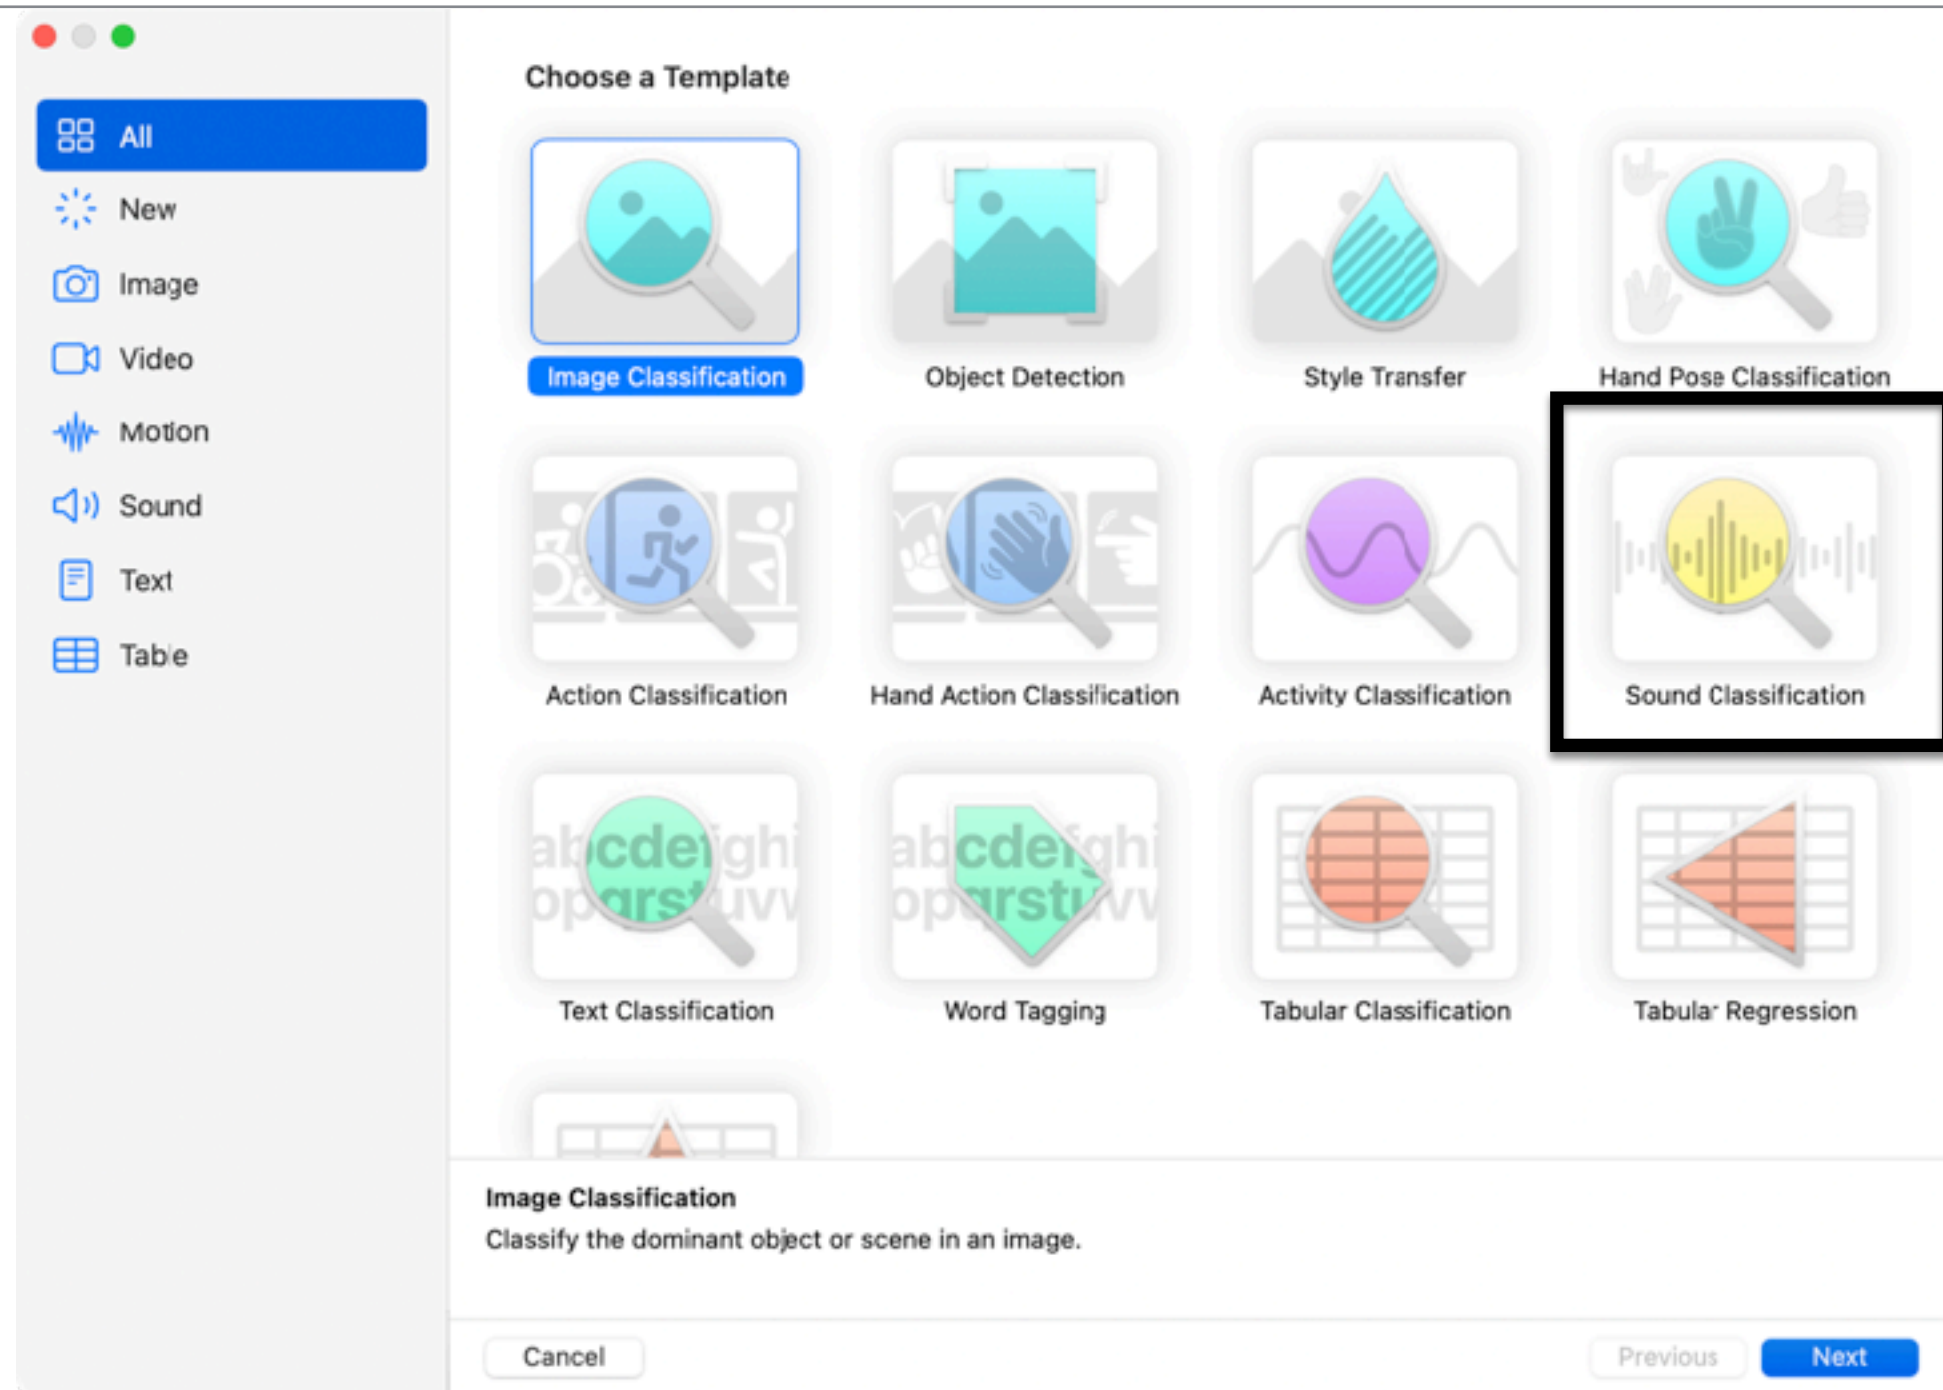
$$\text{Shimmer} = \frac{\Delta 1 - \Delta 2}{\text{mean}}$$

SFSpeechRecognizer

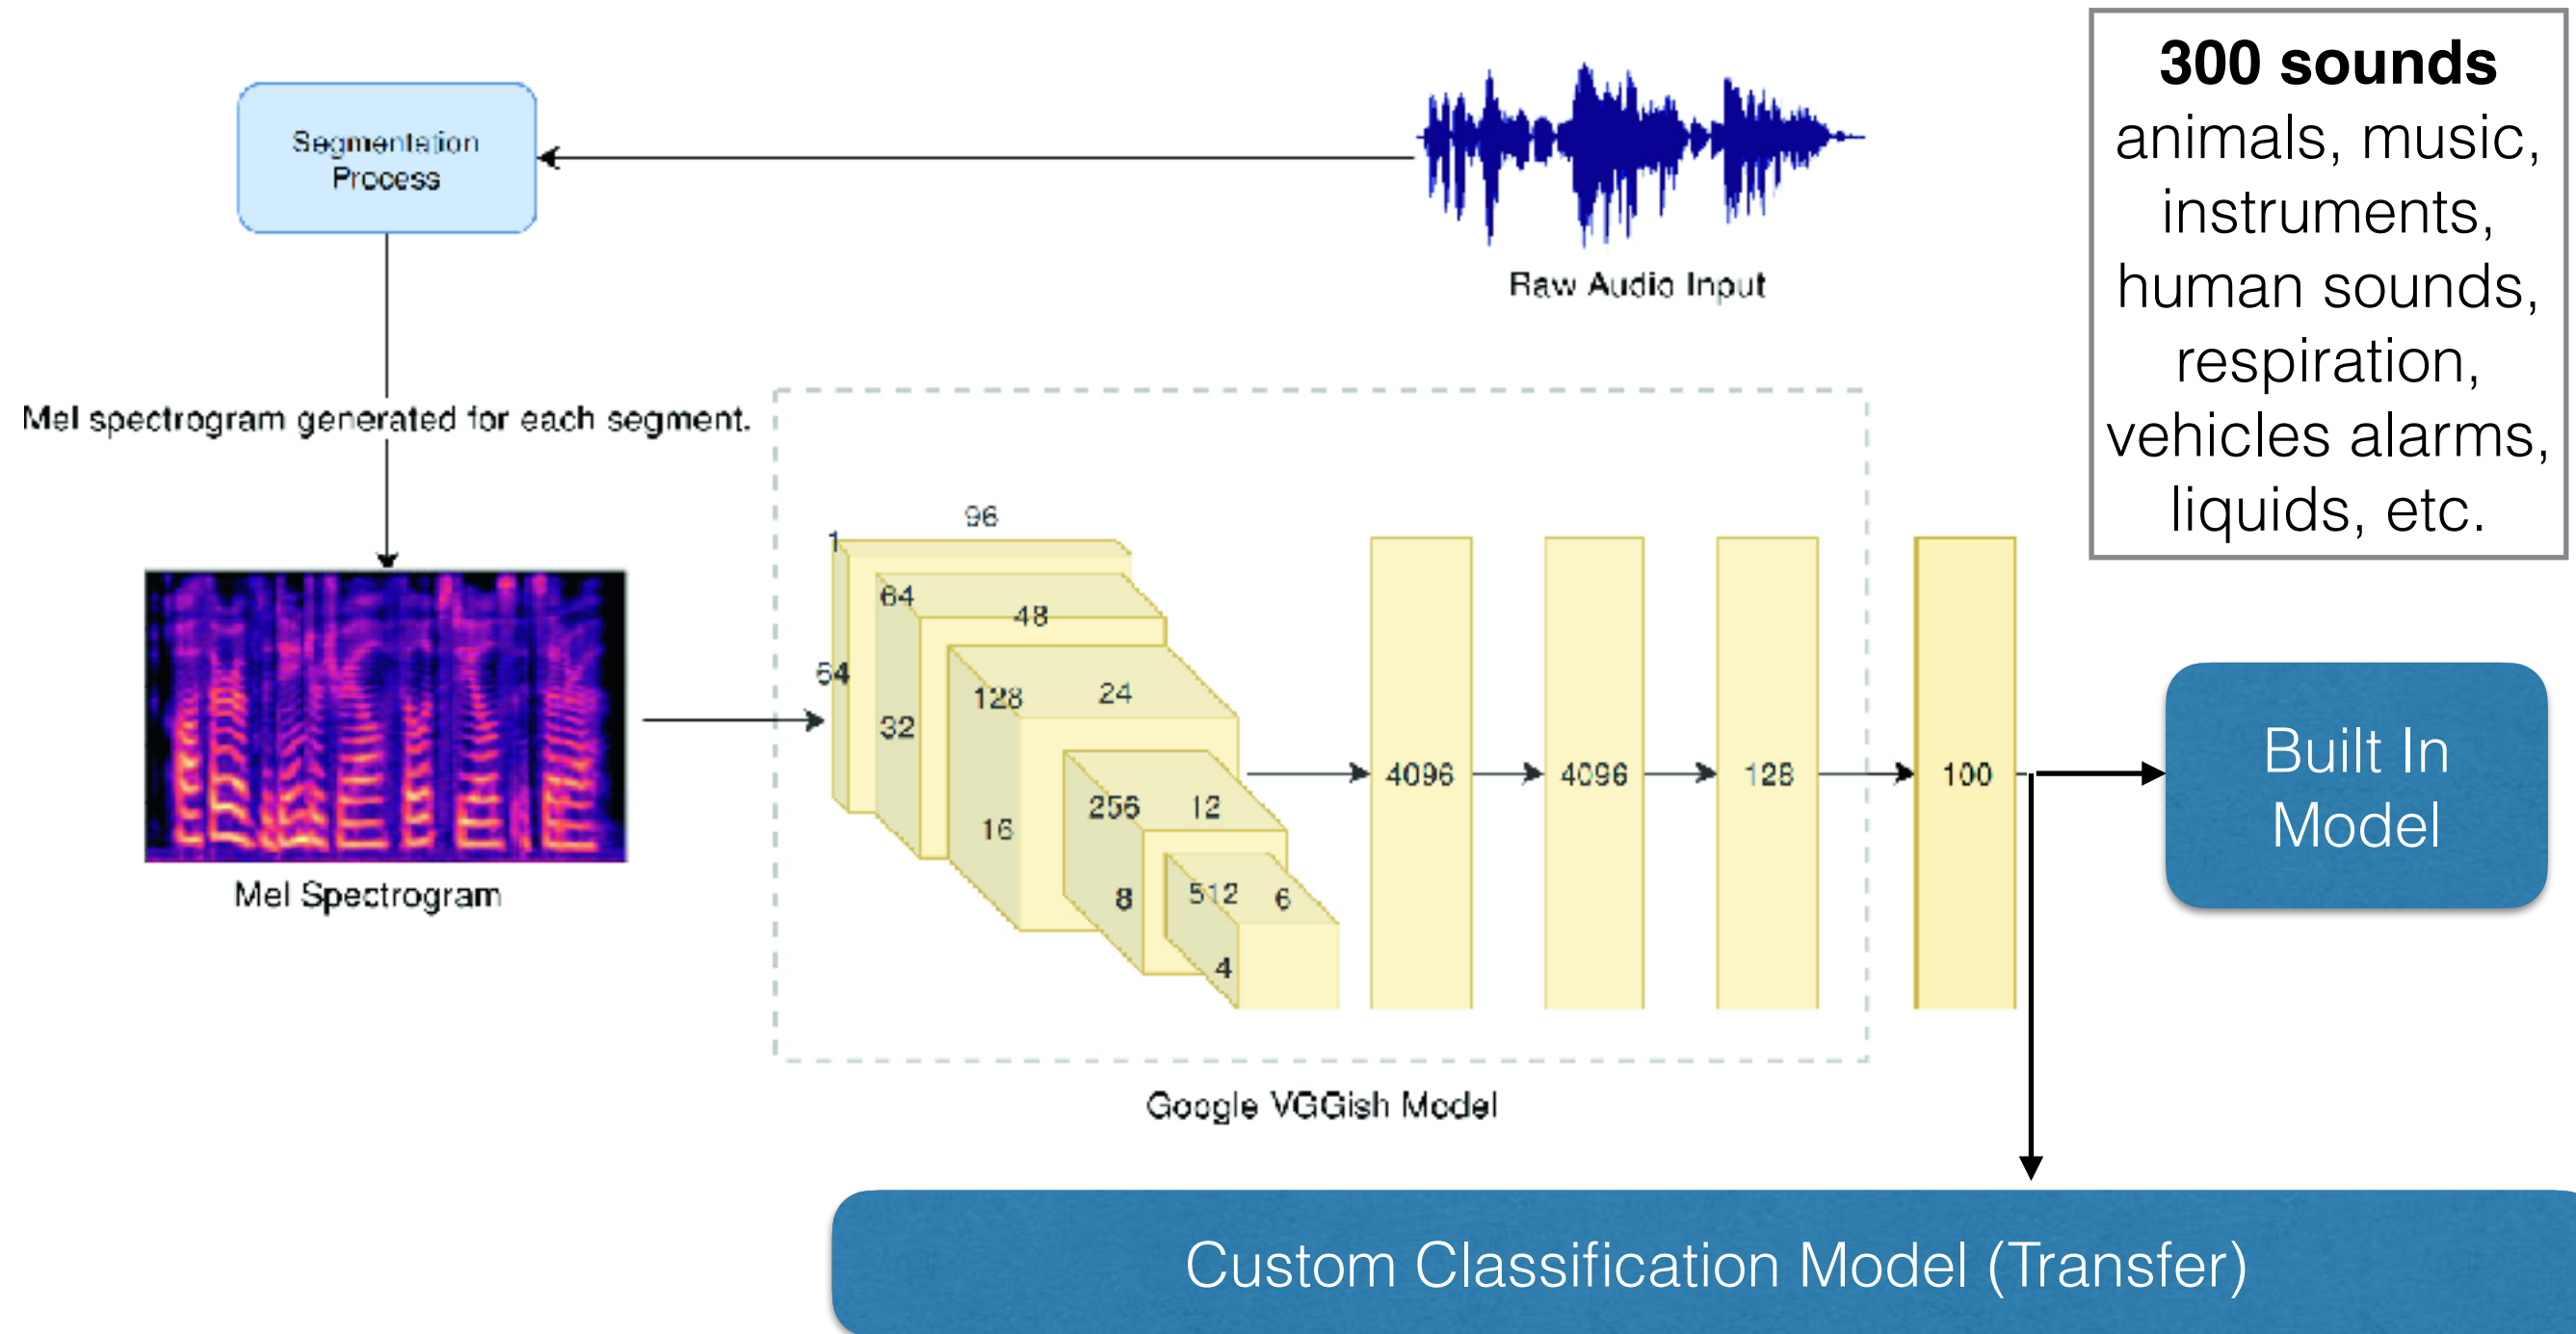
- adding audio blocks from input buffer



Bonus: Create ML (iOS 15.0+)



Create ML Audio Analyzer



Create ML Audio Analyzer

The screenshot displays the 'MySoundClassifier.miproj' app interface. The top bar includes a play button and a 'Share' icon. The main dashboard is divided into four sections: 'Input', 'Metrics', 'Output', and 'Data Inputs'. The 'Input' section shows '11 Classes'. The 'Metrics' section shows '93% Training', '73% Validation', and '39% Testing'. The 'Output' section shows a file icon and '5.5 MB'. The 'Data Inputs' section contains three boxes: 'Training Data' with '6,705 Items' and a dropdown menu showing 'IRMAS-TrainingData'; 'Validation Data' with 'Auto' and a dropdown menu showing 'Automatic'; and 'Testing Data' with '1,432 Items' and a dropdown menu showing 'Part1'. Below these is a 'Parameters' section with 'Maximum Iterations' set to '25' and 'Overlap Factor' set to '50%'. At the bottom, a status bar indicates 'Training completed after 12 minutes, 16 seconds — today at 15:52' and a 'Make a Copy' button.

Section	Value
Input	11 Classes
Metrics	93% Training, 73% Validation, 39% Testing
Output	5.5 MB
Data Inputs	Training Data: 6,705 Items (IRMAS-TrainingData), Validation Data: Auto (Automatic), Testing Data: 1,432 Items (Part1)
Parameters	Maximum Iterations: 25, Overlap Factor: 50%

Training completed after 12 minutes, 16 seconds — today at 15:52

<https://martinmitrevski.com/2019/12/09/sound-classification-with-create-ml-on-ios-13/>

Create ML Audio Analyzer

```
let request = try SNClassifySoundRequest(mlModel: soundClassifier.model)
try streamAnalyzer.add(request, withObserver: self)

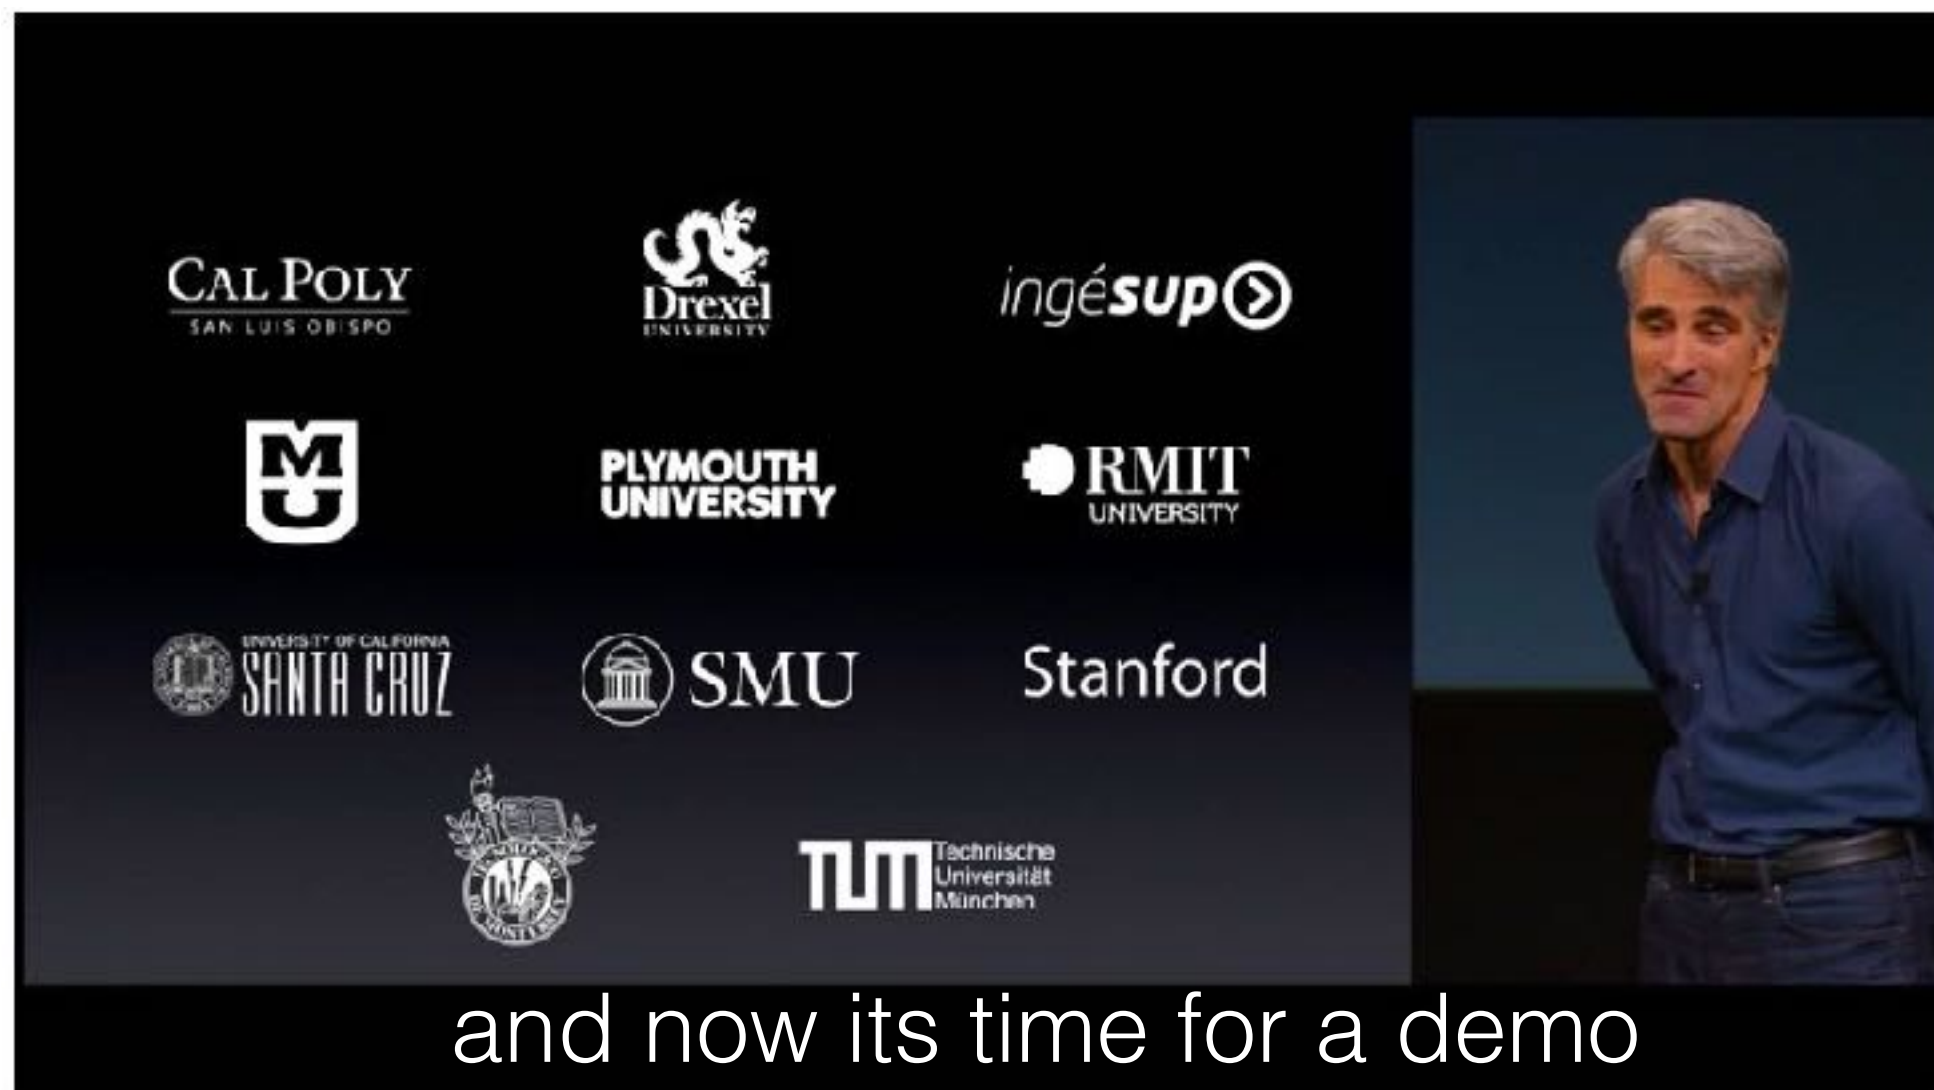
private func prepareForRecording() {
    let inputNode = audioEngine.inputNode
    let recordingFormat = inputNode.outputFormat(forBus: 0)
    streamAnalyzer = SNAudioStreamAnalyzer(format: recordingFormat)
    inputNode.installTap(onBus: 0, bufferSize: 1024, format: recordingFormat) {
        [unowned self] (buffer, when) in
        self.queue.async {
            self.streamAnalyzer.analyze(buffer, atAudioFramePosition: when.sampleTime)
        }
    }
    audioEngine.prepare()
    do { try audioEngine.start() } catch {...}\
}

func request(_ request: SNRequest, didProduce result: SNResult) {
    guard let result = result as? SNClassificationResult else { return }
    for classification in result.classifications {
        print(classification.identifier, classification.confidence)
    }
}
```

<https://martinmitrevski.com/2019/12/09/sound-classification-with-create-ml-on-ios-13/>

Sound Analysis

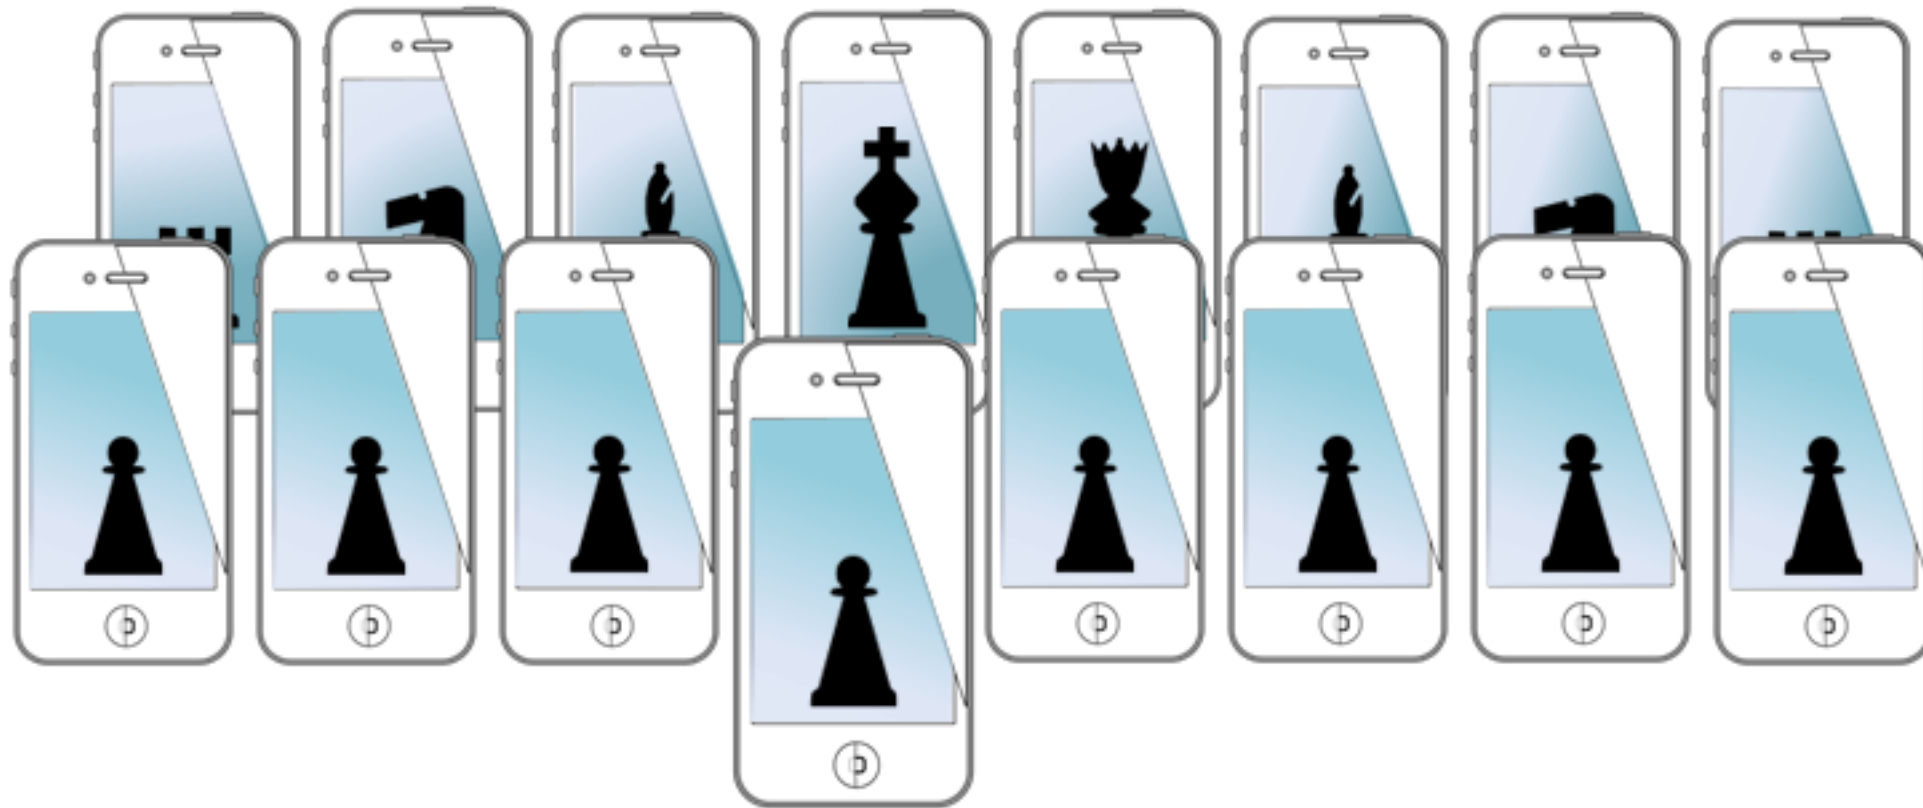
- add sound classification to our project



for next time...

- Pitching
- ~Fin~

MOBILE SENSING LEARNING



CS5323 & 7323

Mobile Sensing and Learning

Speech Recognition and Dictation

Eric C. Larson, Lyle School of Engineering,
Computer Science, Southern Methodist University